

IMAGE RUN:

GENERATED REPORT STATUS: ACTIVE

Diagnostic Laboratory Output

Page 4: Detailed evidence assessment and prediction score disclosures.

DIGITAL EXHIBIT Re-run Scan



INTERACTIVE DIAGNOSTICS

Facial Landmarks

Anomaly Areas

Lighting Angles

Error Level (ELA)

FFT Spectrum

Edge Detection

Wavelet HH

DCT Coefficients

Tip: Hover over markers and red bounding boxes on the media to read deep detailed forensic insights.

INTEGRITY CLASSIFICATION

DEEFAKE DETECTED

HIGH RISK THREAT

Biometric Probability Precision:

99%

ANALYSIS SUMMARY (SIMPLE ENGLISH)

Sightengine classified this media as AI-generated or manipulated with 99% confidence (strongest signal: ai_generated). The forensic indicators support this classification.

TARGET FINDINGS:

Sightengine [type.deepfake] — prob=0.3700, type=deepfake

Sightengine [type.ai_generated] — prob=0.9900, type=ai_generated

[Supporting] Sightengine [type.deepfake] — prob=0.3700, type=deepfake

[Supporting] Sightengine [type.ai_generated] — prob=0.9900, type=ai_generated

[Limitation] Sightengine analysis combines deepfake (face-swap) and genai (fully-generated) model scores from a single API call. Results should be cross-referenced with additional forensic tools.

DOWNLOAD REPORT AS PDF

DEEPPAKE FORENSICS INTELLIGENCE LAB

MILITARY-GRADE COGNITIVE FORENSICS & DIGITAL INTEGRITY VERIFICATION RECORD

HASH_ID: #DF-477

FORENSIC DIAGNOSTIC DECISION: DEEPPAKE ANOMALIES DETECTED

FILE NAME / TARGET ID:	img7_f.jpeg	TIMESTAMP / ACQUIRED:	7/27/2026, 9:02:43 AM
TARGET FORMAT TYPE:	IMAGE	MODEL PROBABILITY:	99% Confidence Score
EVIDENCE CATEGORY:	SUSPECTED MANIPULATION	THREAT SEVERITY:	HIGH LEVEL

I. VECTOR TELEMETRY SIGNATURE & ANOMALY LIST

- [Anomaly Vector 01]: Sightengine [type.deepfake] — prob=0.3700, type=deepfake
- [Anomaly Vector 02]: Sightengine [type.ai_generated] — prob=0.9900, type=ai_generated
- [Anomaly Vector 03]: [Supporting] Sightengine [type.deepfake] — prob=0.3700, type=deepfake
- [Anomaly Vector 04]: [Supporting] Sightengine [type.ai_generated] — prob=0.9900, type=ai_generated
- [Anomaly Vector 05]: [Limitation] Sightengine analysis combines deepfake (face-swap) and genai (fully-generated) model scores from a single API call. Results should be cross-referenced with additional forensic tools.

II. FORENSIC CLINICAL SUMMARY & EXPLANATION

Sightengine classified this media as AI-generated or manipulated with 99% confidence (strongest signal: ai_generated). The forensic indicators support this classification.

III. DISTRIBUTION ANALYSIS & CIRCULATION SUMMARY

Cross-reference with original source if available.

SECURE BLOCKCHAIN HASH ENCRYPTION TAG
VERIFICATION SEAL ID: SHA256-DEEPPAKE-1785123060840

[CERTIFIED DIGITAL EVIDENCE MASTER DECRYPTOR SEAL]

VEDIO RUN:

GENERATED REPORT STATUS: ACTIVE

Diagnostic Laboratory Output

Page 4: Detailed evidence assessment and prediction score disclosures.

DIGITAL EXHIBIT

video_1.mp4

Re-run Scan

0:15 / 0:15

INTERACTIVE DIAGNOSTICS

Facial Landmarks

Anomaly Areas

Lighting Angles

Error Level (ELA)

FFT Spectrum

Edge Detection

Wavelet HH

DCT Coefficients

Tip: Hover over markers and red bounding boxes on the media to read deep detailed forensic insights.

INTEGRITY CLASSIFICATION

VERIFIED GENUINE

LOW RISK THREAT

Biometric Probability Precision:

100%

ANALYSIS SUMMARY (SIMPLE ENGLISH)

Analyzed 6 video frames. Worst-frame AI probability: 0%. Overall: REAL (conf=100%).

TARGET FINDINGS:

Multi-frame analysis: 6 frames analyzed

Worst frame: frame 0 (raw_prob=0.0010)

Frame 0: real (conf=1.00)

Frame 1: real (conf=1.00)

Frame 2: real (conf=1.00)

Frame 3: real (conf=1.00)

Frame 4: real (conf=1.00)

Frame 5: real (conf=1.00)

[Supporting] Frame 0: real (conf=1.00)

[Supporting] Frame 1: real (conf=1.00)

[Supporting] Frame 2: real (conf=1.00)

[Supporting] Frame 3: real (conf=1.00)

[Supporting] Frame 4: real (conf=1.00)

[Supporting] Frame 5: real (conf=1.00)

[Limitation] Video analysis based on extracted key frames, not full frame-by-frame.

DOWNLOAD REPORT AS PDF

DEEPPAKE FORENSICS INTELLIGENCE LAB

MILITARY-GRADE COGNITIVE FORENSICS & DIGITAL INTEGRITY VERIFICATION RECORD

HASH_ID: #DF-482

FORENSIC DIAGNOSTIC DECISION: VERIFIED GENUINE / HIGH INTEGRITY

FILE NAME / TARGET ID:	vid4_r.mp4	TIMESTAMP / ACQUIRED:	7/27/2026, 8:49:04 AM
TARGET FORMAT TYPE:	VIDEO	MODEL PROBABILITY:	100% Confidence Score
EVIDENCE CATEGORY:	AUTHENTIC RECORD	THREAT SEVERITY:	LOW LEVEL

I. VECTOR TELEMETRY SIGNATURE & ANOMALY LIST

- [Anomaly Vector 01]: Multi-frame analysis: 6 frames analyzed
- [Anomaly Vector 02]: Worst frame: frame 0 (raw_prob=0.0010)
- [Anomaly Vector 03]: Frame 0: real (conf=1.00)
- [Anomaly Vector 04]: Frame 1: real (conf=1.00)
- [Anomaly Vector 05]: Frame 2: real (conf=1.00)
- [Anomaly Vector 06]: Frame 3: real (conf=1.00)
- [Anomaly Vector 07]: Frame 4: real (conf=1.00)
- [Anomaly Vector 08]: Frame 5: real (conf=1.00)
- [Anomaly Vector 09]: [Supporting] Frame 0: real (conf=1.00)
- [Anomaly Vector 010]: [Supporting] Frame 1: real (conf=1.00)
- [Anomaly Vector 011]: [Supporting] Frame 2: real (conf=1.00)
- [Anomaly Vector 012]: [Supporting] Frame 3: real (conf=1.00)
- [Anomaly Vector 013]: [Supporting] Frame 4: real (conf=1.00)
- [Anomaly Vector 014]: [Supporting] Frame 5: real (conf=1.00)
- [Anomaly Vector 015]: [Limitation] Video analysis based on extracted key frames, not full frame-by-frame.

II. FORENSIC CLINICAL SUMMARY & EXPLANATION

Analyzed 6 video frames. Worst-frame AI probability: 0%. Overall: REAL (conf=100%).

III. DISTRIBUTION ANALYSIS & CIRCULATION SUMMARY

No further action needed.

SECURE BLOCKCHAIN HASH ENCRYPTION TAG
VERIFICATION SEAL ID: SEAL-DEEPGUARD-1785122245633

[CERTIFIED DIGITAL EVIDENCE MASTER DECRYPTOR SEAL]


```

{
  "status": "success",
  "request": {
    "id": "req_1dW8t5MofnANRqTgQzb",
    "timestamp": 1785122210.116489,
    "operations": 10
  },
  "type": {
    "deepfake": 0.001,
    "ai_generated": 0.001
  },
  "media": {
    "id": "med_1dW8WLPcMTc46kDt3gxwC",
    "uri": "image"
  }
}
INFO:app.providers.sightengine:Candidate 1: Sightengine [type.deepfake] - prob=0.0010, type=deepfake
INFO:app.providers.sightengine:Candidate 2: Sightengine [type.ai_generated] - prob=0.0010, type=ai_generated
INFO:app.providers.sightengine:Selected (max): deepfake
INFO:app.providers.sightengine:Probability: 0.0010
INFO:app.providers.sightengine:Threshold: REAL (FAKE>=0.75, REAL<=0.25)
INFO:app.providers.sightengine:Verdict: real
INFO:app.providers.sightengine:Confidence: 0.9990
INFO:app.providers.sightengine:=====
INFO:app.providers.sightengine:FINAL SCORE from _normalize_verdict: prob=0.0010 type=deepfake verdict=real confidence=0.9990
INFO:app.providers.sightengine:=== SIGHTENGINE NORMALIZED VERDICT ===
{
  "verdict": "real",
  "confidence": 0.999,
  "raw_prob": 0.001,
  "analysis_summary": "Sightengine classified this media as authentic with 100% confidence (max AI/manipulation probability=0.00).",
  "visual_observations": [],
  "forensic_observations": [
    "Sightengine [type.deepfake] \u2014 prob=0.0010, type=deepfake",
    "Sightengine [type.ai_generated] \u2014 prob=0.0010, type=ai_generated"
  ],
  "supporting_evidence": [
    "Sightengine [type.deepfake] \u2014 prob=0.0010, type=deepfake",
    "Sightengine [type.ai_generated] \u2014 prob=0.0010, type=ai_generated"
  ],
  "conflicting_evidence": [],
  "limitations": "Sightengine analysis combines deepfake (face-swap) and genai (fully-generated) model scores from a single API call. Results should be cross-referenced with additional forensic tools.",
  "recommendation": "No further action needed."
}
INFO:app.providers.sightengine:SIGHTENGINE HTTP RESPONSE: status=200 body_preview={
  "status": "success",
  "request": {
    "id": "req_1dW8yXrdQ30Mj5Ab3uNFh",
    "timestamp": 1785122209.397531,
    "operations": 10
  },
  "type": {
    "deepfake": 0.001,
    "ai_generated": 0.001
  },
  "media": {
    "id": "med_1dW8wVd1fzpbu18Jq5zn",
    "uri": "image"
  }
}
INFO:app.providers.sightengine:=== _extract_score TRACE (parser v3) ===
INFO:app.providers.sightengine: Known-key candidate: type=deepfake prob=0.0010 path=type.deepfake
INFO:app.providers.sightengine: Known-key candidate: type=ai_generated prob=0.0010 path=type.ai_generated
INFO:app.providers.sightengine:SELECTED (max): prob=0.0010 type=deepfake path=type.deepfake
INFO:app.providers.sightengine:FINAL SCORE from _extract_score: prob=0.0010 type=deepfake
INFO:app.providers.sightengine:=== _extract_score TRACE END ===
INFO:app.providers.sightengine:===== PARSER TRACE =====
INFO:app.providers.sightengine:Raw JSON:
{
  "status": "success",
  "request": {
    "id": "req_1dW8yXrdQ30Mj5Ab3uNFh",
    "timestamp": 1785122209.397531,
    "operations": 10
  },
  "type": {
    "deepfake": 0.001,
    "ai_generated": 0.001
  },
  "media": {
    "id": "med_1dW8wVd1fzpbu18Jq5zn",
    "uri": "image"
  }
}
INFO:app.providers.sightengine:Candidate 1: Sightengine [type.deepfake] - prob=0.0010, type=deepfake
INFO:app.providers.sightengine:Candidate 2: Sightengine [type.ai_generated] - prob=0.0010, type=ai_generated
INFO:app.providers.sightengine:Selected (max): deepfake
INFO:app.providers.sightengine:Probability: 0.0010
INFO:app.providers.sightengine:Threshold: REAL (FAKE>=0.75, REAL<=0.25)
INFO:app.providers.sightengine:Verdict: real
INFO:app.providers.sightengine:Confidence: 0.9990
INFO:app.providers.sightengine:=====
INFO:app.providers.sightengine:FINAL SCORE from _normalize_verdict: prob=0.0010 type=deepfake verdict=real confidence=0.9990
INFO:app.providers.sightengine:=== SIGHTENGINE NORMALIZED VERDICT ===

```

```

{
  "verdict": "real",
  "confidence": 0.999,
  "raw_prob": 0.001,
  "analysis_summary": "Sightengine classified this media as authentic with 100% confidence (max AI/manipulation probability=0.00).",
  "visual_observations": [],
  "forensic_observations": [
    "Sightengine [type:deepfake] \u2014 prob=0.0010, type=deepfake",
    "Sightengine [type:ai_generated] \u2014 prob=0.0010, type=ai_generated"
  ],
  "supporting_evidence": [
    "Sightengine [type:deepfake] \u2014 prob=0.0010, type=deepfake",
    "Sightengine [type:ai_generated] \u2014 prob=0.0010, type=ai_generated"
  ],
  "conflicting_evidence": [],
  "limitations": "Sightengine analysis combines deepfake (face-swap) and genai (fully-generated) model scores from a single API call. Results should be cross-referenced with additional forensic tools.",
  "recommendation": "No further action needed."
}
INFO:app.providers.sightengine:SIGHTENGINE HTTP RESPONSE: status=200 body_preview={
  "status": "success",
  "request": {
    "id": "req_16W8D6G1X3z2agQlbwInu",
    "timestamp": 1785122210.399374,
    "operations": 10
  },
  "type": {
    "deepfake": 0.001,
    "ai_generated": 0.001
  },
  "media": {
    "id": "med_16W8B75wT8shqW2zHpMuM",
    "uri": "image"
  }
}
INFO:app.providers.sightengine:==== _extract_score TRACE (parser v3) ====
INFO:app.providers.sightengine: Known-key candidate: type=deepfake prob=0.0010 path=type:deepfake
INFO:app.providers.sightengine: Known-key candidate: type=ai_generated prob=0.0010 path=type:ai_generated
INFO:app.providers.sightengine:SELECTED (max): prob=0.0010 type=deepfake path=type:deepfake
INFO:app.providers.sightengine:FINAL SCORE from _extract_score: prob=0.0010 type=deepfake
INFO:app.providers.sightengine:==== _extract_score TRACE END ====
INFO:app.providers.sightengine:===== PARSER TRACE =====
INFO:app.providers.sightengine:Raw JSON:
{
  "status": "success",
  "request": {
    "id": "req_16W8D6G1X3z2agQlbwInu",
    "timestamp": 1785122210.399374,
    "operations": 10
  },
  "type": {
    "deepfake": 0.001,
    "ai_generated": 0.001
  },
  "media": {
    "id": "med_16W8B75wT8shqW2zHpMuM",
    "uri": "image"
  }
}
INFO:app.providers.sightengine:Candidate 1: Sightengine [type:deepfake] - prob=0.0010, type=deepfake
INFO:app.providers.sightengine:Candidate 2: Sightengine [type:ai_generated] - prob=0.0010, type=ai_generated
INFO:app.providers.sightengine:Selected (max): deepfake
INFO:app.providers.sightengine:Probability: 0.0010
INFO:app.providers.sightengine:Threshold: REAL (FAKE)=0.75, REAL<=0.25)
INFO:app.providers.sightengine:Verdict: real
INFO:app.providers.sightengine:Confidence: 0.9990
INFO:app.providers.sightengine:=====
INFO:app.providers.sightengine:FINAL SCORE from _normalize_verdict: prob=0.0010 type=deepfake verdict=real confidence=0.9990
INFO:app.providers.sightengine:===== SIGHTENGINE NORMALIZED VERDICT =====

```

```

{
  "verdict": "real",
  "confidence": 0.999,
  "raw_prob": 0.001,
  "analysis_summary": "Sightengine classified this media as authentic with 100% confidence (max AI/manipulation probability=0.00).",
  "visual_observations": [],
  "forensic_observations": [
    "Sightengine [type:deepfake] \u2014 prob=0.0010, type=deepfake",
    "Sightengine [type:ai_generated] \u2014 prob=0.0010, type=ai_generated"
  ],
  "supporting_evidence": [
    "Sightengine [type:deepfake] \u2014 prob=0.0010, type=deepfake",
    "Sightengine [type:ai_generated] \u2014 prob=0.0010, type=ai_generated"
  ],
  "conflicting_evidence": [],
  "limitations": "Sightengine analysis combines deepfake (face-swap) and genai (fully-generated) model scores from a single API call. Results should be cross-referenced with additional forensic tools.",
  "recommendation": "No further action needed."
}
INFO:app.providers.sightengine:SIGHTENGINE HTTP RESPONSE: status=200 body_preview={
  "status": "success",
  "request": {
    "id": "req_16W8B8CV4fHaLR6LhVAlz",
    "timestamp": 1785122209.395549,
    "operations": 10
  },
  "type": {
    "deepfake": 0.001,
    "ai_generated": 0.001
  },
  "media": {
    "id": "med_16W8Y0VYF0ZxjybhXeMw",
    "uri": "image"
  }
}
INFO:app.providers.sightengine:==== _extract_score TRACE (parser v3) ====
INFO:app.providers.sightengine: Known-key candidate: type=deepfake prob=0.0010 path=type:deepfake
INFO:app.providers.sightengine: Known-key candidate: type=ai_generated prob=0.0010 path=type:ai_generated
INFO:app.providers.sightengine:SELECTED (max): prob=0.0010 type=deepfake path=type:deepfake
INFO:app.providers.sightengine:FINAL SCORE from _extract_score: prob=0.0010 type=deepfake
INFO:app.providers.sightengine:==== _extract_score TRACE END ====
INFO:app.providers.sightengine:===== PARSER TRACE =====
INFO:app.providers.sightengine:Raw JSON:
{
  "status": "success",
  "request": {
    "id": "req_16W8B8CV4fHaLR6LhVAlz",
    "timestamp": 1785122209.395549,
    "operations": 10
  },
  "type": {
    "deepfake": 0.001,
    "ai_generated": 0.001
  },
  "media": {
    "id": "med_16W8Y0VYF0ZxjybhXeMw",
    "uri": "image"
  }
}
INFO:app.providers.sightengine:Candidate 1: Sightengine [type:deepfake] - prob=0.0010, type=deepfake
INFO:app.providers.sightengine:Candidate 2: Sightengine [type:ai_generated] - prob=0.0010, type=ai_generated
INFO:app.providers.sightengine:Selected (max): deepfake
INFO:app.providers.sightengine:Probability: 0.0010
INFO:app.providers.sightengine:Threshold: REAL (FAKE)=0.75, REAL<=0.25)
INFO:app.providers.sightengine:Verdict: real
INFO:app.providers.sightengine:Confidence: 0.9990
INFO:app.providers.sightengine:=====
INFO:app.providers.sightengine:FINAL SCORE from _normalize_verdict: prob=0.0010 type=deepfake verdict=real confidence=0.9990
INFO:app.providers.sightengine:===== SIGHTENGINE NORMALIZED VERDICT =====

```

```

{
  "verdict": "real",
  "confidence": 0.999,
  "raw_prob": 0.001,
  "analysis_summary": "Sightengine classified this media as authentic with 100% confidence (max AI/manipulation probability=0.00).",
  "visual_observations": [],
  "forensic_observations": [
    "Sightengine [type:deepfake] \u2014 prob=0.0010, type=deepfake",
    "Sightengine [type:ai-generated] \u2014 prob=0.0010, type=ai-generated"
  ],
  "supporting_evidence": [
    "Sightengine [type:deepfake] \u2014 prob=0.0010, type=deepfake",
    "Sightengine [type:ai-generated] \u2014 prob=0.0010, type=ai-generated"
  ],
  "conflicting_evidence": [],
  "limitations": "Sightengine analysis combines deepfake (face-swap) and genai (fully-generated) model scores from a single API call. Results should be cross-referenced with additional forensic tools.",
  "recommendation": "No further action needed."
}
INFO:app.providers.sightengine: SIGHTENGINE HTTP RESPONSE: status=200 body_preview=[
  "status": "success",
  "request": {
    "id": "req_16W8MqfMAwZPLXgey6Iw",
    "timestamp": 1785122210.136482,
    "operations": 10
  },
  "type": {
    "deepfake": 0.001,
    "ai_generated": 0.001
  },
  "media": {
    "id": "med_16W8TvzH4Hw77lvc2Mbvxx",
    "uri": "image"
  }
]
INFO:app.providers.sightengine:==== _extract_score TRACE (parser v3) ====
INFO:app.providers.sightengine: Known-key candidate: type=deepfake prob=0.0010 path=type:deepfake
INFO:app.providers.sightengine: Known-key candidate: type=ai-generated prob=0.0010 path=type:ai-generated
INFO:app.providers.sightengine:SELECTED (max): prob=0.0010 type=deepfake path=type:deepfake
INFO:app.providers.sightengine:FINAL SCORE from _extract_score: prob=0.0010 type=deepfake
INFO:app.providers.sightengine:==== _extract_score TRACE END ====
INFO:app.providers.sightengine:===== PARSE TRACE =====
INFO:app.providers.sightengine:Raw JSON:
{
  "status": "success",
  "request": {
    "id": "req_16W8MqfMAwZPLXgey6Iw",
    "timestamp": 1785122210.136482,
    "operations": 10
  },
  "type": {
    "deepfake": 0.001,
    "ai_generated": 0.001
  },
  "media": {
    "id": "med_16W8TvzH4Hw77lvc2Mbvxx",
    "uri": "image"
  }
}
INFO:app.providers.sightengine:Candidate 1: Sightengine [type:deepfake] - prob=0.0010, type=deepfake
INFO:app.providers.sightengine:Candidate 2: Sightengine [type:ai-generated] - prob=0.0010, type=ai-generated
INFO:app.providers.sightengine:Selected (max): deepfake
INFO:app.providers.sightengine:Probability: 0.0010
INFO:app.providers.sightengine:Threshold: REAL (FAKE)>=0.75, REAL<=0.25)
INFO:app.providers.sightengine:Verdict: real
INFO:app.providers.sightengine:Confidence: 0.9990
INFO:app.providers.sightengine:=====
INFO:app.providers.sightengine:FINAL SCORE from _normalize_verdict: prob=0.0010 type=deepfake verdict=real confidence=0.9990
INFO:app.providers.sightengine:===== SIGHTENGINE NORMALIZED VERDICT =====

```

```

{
  "verdict": "real",
  "confidence": 0.999,
  "raw_prob": 0.001,
  "analysis_summary": "Sightengine classified this media as authentic with 100% confidence (max AI/manipulation probability=0.00).",
  "visual_observations": [],
  "forensic_observations": [
    "Sightengine [type:deepfake] \u2014 prob=0.0010, type=deepfake",
    "Sightengine [type:ai-generated] \u2014 prob=0.0010, type=ai-generated"
  ],
  "supporting_evidence": [
    "Sightengine [type:deepfake] \u2014 prob=0.0010, type=deepfake",
    "Sightengine [type:ai-generated] \u2014 prob=0.0010, type=ai-generated"
  ],
  "conflicting_evidence": [],
  "limitations": "Sightengine analysis combines deepfake (face-swap) and genai (fully-generated) model scores from a single API call. Results should be cross-referenced with additional forensic tools.",
  "recommendation": "No further action needed."
}
INFO:app.runner:Sightengine frame 1/6: verdict=real conf=0.9990 raw_prob=0.0010
INFO:app.runner:Sightengine frame 2/6: verdict=real conf=0.9990 raw_prob=0.0010
INFO:app.runner:Sightengine frame 3/6: verdict=real conf=0.9990 raw_prob=0.0010
INFO:app.runner:Sightengine frame 4/6: verdict=real conf=0.9990 raw_prob=0.0010
INFO:app.runner:Sightengine frame 5/6: verdict=real conf=0.9990 raw_prob=0.0010
INFO:app.runner:Sightengine frame 6/6: verdict=real conf=0.9990 raw_prob=0.0010
INFO:app.runner:Sightengine multi-frame: frames=6 max_raw_prob=0.0010 worst_frame=0 verdict=real confidence=0.9990
INFO:app.runner:SIGHTENGINE RAW RESPONSE: {
  "verdict": "real",
  "confidence": 0.999,
  "raw_prob": 0.001,
  "frame_analysis": [
    {
      "frame_index": 0,
      "verdict": "real",
      "confidence": 0.999,
      "raw_prob": 0.001,
      "summary": "Sightengine classified this media as authentic with 100% confidence (max AI/manipulation probability=0.00).",
    },
    {
      "frame_index": 1,
      "verdict": "real",
      "confidence": 0.999,
      "raw_prob": 0.001,
      "summary": "Sightengine classified this media as authentic with 100% confidence (max AI/manipulation probability=0.00).",
    },
    {
      "frame_index": 2,
      "verdict": "real",
      "confidence": 0.999,
      "raw_prob": 0.001,
      "summary": "Sightengine classified this media as authentic with 100% confidence (max AI/manipulation probability=0.00).",
    },
    {
      "frame_index": 3,
      "verdict": "real",
      "confidence": 0.999,
      "raw_prob": 0.001,
      "summary": "Sightengine classified this media as authentic with 100% confidence (max AI/manipulation probability=0.00).",
    },
    {
      "frame_index": 4,
      "verdict": "real",
      "confidence": 0.999,
      "raw_prob": 0.001,
      "summary": "Sightengine classified this media as authentic with 100% confidence (max AI/manipulation probability=0.00).",
    },
    {
      "frame_index": 5,
      "verdict": "real",
      "confidence": 0.999,
      "raw_prob": 0.001,
      "summary": "Sightengine classified this media as authentic with 100% confidence (max AI/manipulation probability=0.00).",
    }
  ]
}

```

```

"analysis_summary": "Analyzed 6 video frames. Worst-frame AI probability: 0%. Overall: REAL (conf=100%).",
"visual_observations": [],
"forensic_observations": [
  "Multi-frame analysis: 6 frames analyzed",
  "Worst frame: frame 0 (raw_prob=0.0010)",
  "Frame 0: real (conf=1.00)",
  "Frame 1: real (conf=1.00)",
  "Frame 2: real (conf=1.00)",
  "Frame 3: real (conf=1.00)",
  "Frame 4: real (conf=1.00)",
  "Frame 5: real (conf=1.00)"
],
"supporting_evidence": [
  "Frame 0: real (conf=1.00)",
  "Frame 1: real (conf=1.00)",
  "Frame 2: real (conf=1.00)",
  "Frame 3: real (conf=1.00)",
  "Frame 4: real (conf=1.00)",
  "Frame 5: real (conf=1.00)"
],
"conflicting_evidence": [],
"limitations": "Video analysis based on extracted key frames, not full frame-by-frame.",
"recommendation": "No further action needed."
}
INFO:app.runner:Sightengine succeeded: verdict=real confidence=1.00 frames=6 - prepending result for LLM reconciliation
INFO:app.runner:Sightengine clear verdict (conf=0.9990 >= 0.8) - returning directly, no LLM reconciliation
INFO:app.runner:=== PIPELINE: Analysis stage complete ===
INFO:app.runner: incoming verdict=real confidence=0.9990 model=sightengine fallback=False
INFO:app.runner:RUNNING: D:\adk-workspace\deepguard-ai\backend\app\runner.py
INFO:app.runner:Function: _run_report_with_fallback()
INFO:app.runner:Trying Report Primary: Groq (groq/llama-3.3-70b-versatile)
WARNING:app.runner:Report Primary (Groq) FAILED
INFO:app.runner:Trying Report Fallback 1: NVIDIA Omni (nvidia_nim/nvidia/nemotron-3-nano-omni-30b-a3b-reasoning)
INFO:app.runner:Report Fallback 1 (NVIDIA) succeeded: 6055 chars
INFO:app.runner:=== PIPELINE: Report stage complete ===
INFO:app.runner: model=nvidia_nim/nvidia/nemotron-3-nano-omni-30b-a3b-reasoning degraded=True
INFO:app.runner:=== PIPELINE: Entering output assembly ===
INFO:app.runner: incoming verdict=real confidence=0.9990
INFO:app.runner:=== VERDICT CHAIN ===
INFO:app.runner: Provider confidence: 0.9990
INFO:app.runner: Final verdict: real (conf=0.9990)
INFO:app.runner:Report verdict: real (conf=0.9990)
INFO:app.runner:PDF generated: 6197 bytes
INFO:app.services.audit_service:Audit entry written - request=7fcd3a17f9184ee0 verdict=real hash=54ed7514998f
INFO:app.runner:API verdict: real (conf=0.9990)
INFO:app.runner:=== VERDICT CHAIN COMPLETE ===
INFO: 127.0.0.1:58101 - "POST /api/analyze HTTP/1.1" 200 OK
INFO:watchfiles.main:2 changes detected
|

```

